

Interviewing in Ten Minutes

Spencer Tipping

November 25, 2014

Contents

1	The business protocol	2
2	The interviewer	3
2.1	Aside: interview attire	3
3	The résumé and cover letter	4
4	The technical phone interview	5
4.1	Technical literacy	5
4.2	Development literacy	6
4.3	Career/self literacy	6
5	Coding interview	7
5.1	Coding challenge etiquette	8
5.2	Quick-fail problems	9
5.2.1	FizzBuzz	9
5.2.2	Sum of linked list	11
5.2.3	Practice problems	12
5.3	Algorithm problems	12
5.3.1	Maximum subarray	12
5.3.2	Parse simple math	15
5.3.3	Sorted word index	16
5.4	High-level design questions	18
5.4.1	IO scheduler	19
5.5	Practice problem solutions	21

About this guide

Over the past few years I've worked for a lot of startups. (Because I job-hopped, not because I'm important.) For one reason or another, I ended up conducting technical interviews at about half of these companies. Usually we were looking

for developers who knew some high-end functional language, so this guide is strongly biased by that – though I think it also applies to more standard startup jobs.

Everything in this guide is my own opinion. It does not reflect the views of any of my employers, past or present.

1 The business protocol

For various reasons, partly tradition and partly expedience, companies tend to have a set of norms that dictate how people work together. From what I've been able to observe, the following commandments encode the business protocol:

1. **Thou shalt respect other people's time like thine own.** This is absolutely crucial. Business is about money, people convert money to time, and if you waste/take an employee's time then you're wasting the company's money. People holding the money to pay your salary want to see a positive return on their investment in you, which is undermined if you then make other people less productive. **Be the person who makes others more productive.**
2. **Thou shalt get things done.** By this I mean, if someone gives you a task it will be done. They don't have to think twice about it. If you hit difficulty, you'll figure it out. If you can prove it's impossible, change the problem to be possible and explain why you did. **Be the person who gives up last.**
3. **Thou shalt work only when it is productive.** Work \neq productivity. The two are nearly orthogonal, in fact. Understand what is productive at as high a level as you can, and work on that stuff. Busywork is worse than doing nothing because you'll be unavailable but will be getting nothing done. **Be the person who does everything for a reason.**
4. **Thou shalt simplify.** This is a way to respect people's time. Every problem that comes your way should be simpler after it interacts with you. Make things look easy because to you they are easy. **Be the person who finds the best solution.**
5. **Thou shalt be fluent.** This goes two ways: you must communicate flawlessly both with people and with computers. Functionally, a programmer is an information channel that translates between human intent and computational action. You want this channel to be efficient, lossless, and high-capacity. **Be the best conduit of intent from people to machines.**
6. **Thou shalt not bikeshed.** Thou shalt not argue in general, in fact. But if you must, make sure it's about an issue that justifies a conversation costing the company more than \$100/hour (nearly every conversation is likely to cost this, and some will cost far more). **Be the person who cuts through the noise and focuses on real problems.**

I'll go into more details about how to convey these things during an interview later on. For now, though, understand that the above values (or similar ones) serve as the conduit by which your skills can be applied to the company's problems. If this conduit is missing or sufficiently damaged, your skills won't matter.

2 The interviewer

Interviewers are not usually sadists. While the case could be made that the job demands a certain degree of malevolence, most of the confrontation during the interview is due to the economics underlying hiring.

Consider the common flow of resume screen, phone screen, and then in-person interviews. There's a progression of risk/cost associated with moving forward at each stage: screening a resume requires minutes, a phone interview requires an hour, and an in-person interview requires several people-hours, possibly with hotel and airfare. In dollar amounts, these could easily represent investments of \$10, \$100, and \$1000.¹ The worst case, a bad hire, costs potentially \$10,000 or more.

Given this, the interviewer operates as an investor of the company's assets.² Their job is to evaluate risks for the company by making bets on people. Because of this, if you are an interviewee, your job is to sell such a bet to the person interviewing you. Perhaps more to the point, the interviewer is likely to be conservative because they don't get paid commission; so your goal is to be so awesome that they feel safer saying yes than no.

2.1 Aside: interview attire

The interviewer does not care what you wear to an interview. Nobody cares about this except you. Unlike in almost any other industry, a suit is much more likely to be a liability than anything else. I have gotten every job I've had by going to interviews in flip-flops, khakis, and a T-shirt, and being very comfortable because this is what I'm used to wearing. The interview is about your mind, so give your mind the least-distracting atmosphere possible so it can do its job.

It makes perfect sense if you think about the interviewer's position. If you're awesome, the interviewer cannot go back to the hiring manager with an assessment like, "well, the candidate is the most productive we've ever seen, but can we really put up with someone who dresses that way?" "Fashion-aware" and "software developer" don't even belong in the same sentence. The

¹I've been told that in-person interviews can cost significantly upwards of \$1000, so take these numbers with a grain of salt and treat them as lower bounds.

²Unlike most investors, however, their pay is constant; so they don't have a stake in the outcome of the interview. This creates a slight negative bias, but as far as I know it shouldn't influence their decisions otherwise.

interviewer will probably say something like, “the candidate blew us away on the coding tests, *and this one even wore pants!*” “Capable *and* pants? Hire.”

The flip side is equally clear. If you’re awkward and uncomfortable, the interviewer doesn’t know whether it’s because you’re having trouble with the problems, or whether it’s that you’re wearing a suit to an office with a beer pong table and a CEO dressed in a hoodie. The interviewer is certainly not going to say, “well, that candidate can’t code his way out of paper bag but man is he a sharp dresser. I’m a yes.”

3 The résumé and cover letter

At \$100, a phone screen is still fairly cheap. The stakes of a resume and cover letter are not usually very high, nor is there very much signal for the interviewer to work with. However, if you’re about optimizing stuff (and what engineer isn’t?), there are a few things you can do to communicate at this point:

1. **T_EX** your resume. There’s only upside here, and it’s a nice positive signal that you know what you’re doing. If you don’t know T_EX, use HTML and export as PDF. If all else fails, write one in plain text and claim that minimalism is a social virtue you take seriously. Don’t use Word because it conveys that you have no standards. As a colleague of mine once said, “if I’m on the fence and the resume is a word doc, I typically reject the candidate.” **Use a spellchecker.**
2. Keep your resume down to one page. This shows that you respect their time. **Use a spellchecker.**
3. Write a purposeful but informal cover letter. Be personable, reasonably enthusiastic, and talk about how the stuff you do for fun is relevant to the company’s hard problems. **Use a spellchecker.**
4. Link to your github profile. If you don’t have one yet, create one and upload something. If you don’t have anything to upload, start solving Project Euler and upload that. **Use a spellchecker.**
5. Put some code in your cover letter. Nothing elaborate, but just something you thought was cool, or maybe a challenge problem they have on their website, or something. Do this only if it makes sense, but it’s a great opportunity if you can. **Use a spellchecker.**
6. Link to a project you worked on that really describes who you are. **Use a spellchecker.**

There are a few things you absolutely cannot do:

1. Talk about salary in any way. It’s just too early and it sends the wrong message. It would be like someone asking about your finances before the first date. This is a question you ask when you’re considering marriage, not dinner.

2. Lie. If you're called out on it, you're done for and people will hate you. Word might get around to other companies if you're unlucky. If the truth isn't good enough, fix the truth.
3. Be noticeably unprofessional. Startups aren't overwhelmingly formal places, but they're just as vulnerable to harassment lawsuits as anyone else. You do not want to give anyone reason to believe that you're going to cause legal problems or alienate people.
4. Bring politics or religion to the table. These topics are emotional catalysts with no upside and enormous potential for problems. Brendan Eich, the inventor of Javascript, was forced to resign as CEO of Mozilla because he had made a \$1000 contribution against gay marriage six years prior. If you have public political statements on a blog or elsewhere, take them down.

4 The technical phone interview

At this point you have cost the company \$110 or so, and they're deciding whether or not to bump that figure by a factor of 10. This is real money, so the burden is on you to convince the interviewer (who may or may not be good at interviewing people) that you're the best bet they'll make this year. Depending on the job you're applying for, there are a few different ways to do that.

In this section I'm going to assume the technical interview is just talking on the phone. The next section is entirely devoted to how to handle a collaborative coding interview.

4.1 Technical literacy

This is an absolute must. At a minimum, you need to be able to answer questions like these with little or no thought:

- Write the FizzBuzz function.
- Reverse an array in place.
- Reverse a linked list in linear time.
- What is the insertion time for a binary tree?
- What is the insertion time for a hashtable (worst, average)?
- How is a priority queue implemented?
- Quicksort vs mergesort, pros/cons?
- What does $O(n^2)$ mean formally?
- What does "atomic" mean?

- What is memoization?
- Why do floating-point computations often use an epsilon?
- How are negative integers represented?
- Write the Fibonacci function without using recursion.
- What's the difference between the stack and the heap?
- Conceptually, what does a garbage collector do?
- Conceptually, how does virtual memory work?
- What happens under the hood when you go to `google.com`?
- Why should you profile before optimizing?
- What is an example use case for a cache?
- How would you do [some simple task] in a UNIX shell?

4.2 Development literacy

This is also important, especially if you're applying for a position that requires experience. You should be able to speak intelligently about things like:

- The bugs you most commonly write, and habits you've developed to avoid them.
- How expensive it is to fix a bug in production.
- Strengths and weaknesses of unit testing.
- The drawbacks of using powerful languages like Lisp. (You can extol their virtues too, but if that's all you do then people are unlikely to take you very seriously.)
- Your preferences about design and documentation, and why you do it this way.
- A project you've worked on, how it turned out, and how your actions contributed to that outcome.

4.3 Career/self literacy

This is less important from a technical perspective, but it still sends a signal and is worth getting right. You should have a reasonably well-informed opinion about:

- Your preferred development style/environment

- How you like to interact with your manager
- How autonomous you prefer to be (more is not always better)
- The types of projects you like to work on, and why
- Roles that don't suit you
- How you envision you might be useful to this company
- How you want to change over the coming months/years
- What it takes to keep you at a job
- What would make you leave a job
- What would make you leave a job and write a blog post about why you did

This is the part where you are interviewing the company. Just because they offer you a job doesn't mean you should take it. Taking the wrong job is not only bad for personal growth, but will most likely lead to a minefield of a resume that will hurt your future career prospects. Demonstrating reliability and the ability to stick it out is important in the long run for two reasons. One is that hiring is expensive, so having to replace you quickly is a liability. But the other is that people who tend to job-hop might be doing it because they're bad at working with people. You want to eliminate this possibility in the mind of your interviewer.

5 Coding interview

People tend to assume there are two outcomes for a coding interview: succeed or fail. This is not technically true, however; a third outcome is *succeed specatularly*, which is what you should do if you have the option, and here's why.

Coding challenge interviewers are usually engineers, and engineers tend to like to have answers to questions so they can move on. Because of this, they'll try to jump to a conclusion about you as soon as they have enough signal to justify it. This conclusion is rarely final, but at least they have something to work with if they don't observe anything else.

The interview question will usually fall into one of two categories, each equally expedient. One is stuff like FizzBuzz, which rapidly detects certain failure modes. The other is something open-ended that gives you some room to be creative. Although these problems can detect failure, they're more often about detecting excellence.

If at any point in the interview process someone decides you're excellent, you have a much higher chance of getting an offer. Picasso was excellent and his paintings are now worth millions, whereas the guy who paints imitation art just like Picasso is broke. Setting a new bar instantly promotes you into the

class of people who can't be replaced, and this makes it much easier to justify further investment in interviewing you.

5.1 Coding challenge etiquette

Enjoy the coding challenge. Relish it. This is an opportunity to be asked difficult, interesting questions that are designed to be solvable in less than an hour. It's break from all the mundane aspects of writing production code, and an opportunity to prove to yourself as well as your interviewer that you can solve anything that comes your way.

- Have a language you're really comfortable with. Java is not a bad choice for this because it's straightforward and everyone knows it.
- Unless the problem is trivial, verbally sketch out your thought process. Argue against the approaches you've taken to see if they hold up. Don't write any code until you know what you're doing. The goal of all this is to demonstrate that your thought process is effective at solving problems. Silence followed by perfection also works, but is riskier.
- Make simplifying assumptions as necessary, but be prepared to undo them later. Always mention when you're doing this.
- If the problem is unsolvable as stated, change it minimally to make it solvable and explain why you need to do this.
- If the problem is too vague or ill-specified to even work with, talk with the interviewer to work the problem into something you can solve. Everyone loses if you walk away from the interview without having done anything.
- When writing code, use idioms that reflect a considered perspective. It should be simple, easy to maintain, well-suited to the problem, and generally self-documenting. You should be able to intelligently discuss every aspect of your coding style.
- Never sweat the small stuff. Nobody cares if you know what the deletion method is called on Java's `HashSet` class. Just casually mention that you'd look it up.
- Never get defensive. If you've made a mistake, fix it quickly (or think about it and discuss it) and move on. Any good interviewer will understand that people make mistakes, and some will even set up questions specifically designed to mislead you. This is all part of the interview and it's to see how you take it when things go wrong.

5.2 Quick-fail problems

These are the easiest problems you'll get, so you need to make them look easy. Some classic examples:

1. Write the factorial function.
2. Write FizzBuzz.
3. Calculate the sum of numbers in a linked list.
4. Print the last line of a file (naive solution).

Coding style matters, as well as speed and whether you can write them without any bugs. I'll go through some solutions of varying quality.

5.2.1 FizzBuzz

Bad solutions indicate a deep lack of understanding of some sort. Any competent interviewer will count a solution like this against you:

Listing 1 examples/fizzbuzz-bad.java

```
1 public int fizz_buzz() {           // "public" -- eh; "int" ... ????
2     int count = 0;                 // count of what?; also, indent!
3     while(true) {                  // forever? why?
4         int count2 = count          // "count2"?
5         while (count2 > 0) {         // modulus! modulus!
6             count2 = count2 - 3;
7         }
8         int count3 = count;          // ok, this is just silly
9         while (count3 > 0) {
10            count3 = count3 - 5;
11        }
12        //check for divisibility
13        if (count2 == 0){
14            if ( count3 == 0) {
15                System.out.println("fizzbuzz");
16            }
17            else                      // missing open brace, not good
18                System.out.println ("fizz");
19        }
20        else
21            if count3 == 0 {
22                System.out.println("buzz")
23            }
24        else {
25            System.out.println(count)
26        }
```

```
27 }
28 }
```

This, on the other hand, is a solution I would consider to be nearly perfect:

Listing 2 examples/fizzbuzz-good.java

```
1 for (int i = 1; i < 100; ++i) {
2     final boolean by3 = i % 3 == 0;
3     final boolean by5 = i % 5 == 0;
4     if (by3 && by5)
5         println("fizzbuzz");
6     else if (by5)
7         println("buzz");
8     else if (by3)
9         println("fizz");
10    else
11        println(i);
12 }
```

Some important differences:

- Indentation and good style should be habit for you. Code should have meaning, not just behavior. Formatting should be absolutely consistent (any random-looking inconsistency such as in the bad solution above will lead your interviewer to suspect you might be junior at writing code, and you really want to avoid this).
- Know the nuances and idioms of your language. Java gives you a modulus operator, so you should use it here. Use the appropriate iteration and decisional constructs.
- Draw attention to stuff that matters. Here, most of the space is taken up by the condition structure, which is parameterized by two variables calculated above. The `for` loop is appropriately minimal because it's a standard iteration form that uses a generic variable name.
- Write the solution to resemble the problem. Here, `by3` and `by5` are the two conditions we care about. This makes the `if` stack nearly a verbatim translation of the problem statement. Because of this, it's very easy to look at the code and be confident it does the right thing. These variables are short, descriptive, and not generic, which causes them to draw attention without taking up much space.
- Write hints into your code (as appropriate). I like to use `final` to indicate invariance, for example. This is a matter of opinion, but anything you can do to make it clear that you're thinking about maintainability will work in your favor.

5.2.2 Sum of linked list

This problem is nice because it has a mediocre solution:

Listing 3 examples/sumlist-mediocre.java

```
1 interface OurList {
2     int head();
3     OurList tail();
4 }
5 int sum(OurList list) {
6     if (list == null)
7         return 0;
8     else
9         return list.head() + sum(list.tail());
10 }
```

This is logically correct, terse, and generally well-expressed. But it has two big problems. First, it returns an `int` rather than a `long` or `double`; this is probably a mistake. Second, and more importantly, it requires linear stack space and may cause an overflow. Java does not optimize tail calls, but even if it did the recursive call here is not in tail position. A much better solution uses iteration:

Listing 4 examples/sumlist-good.java

```
1 interface OurList {
2     int head();
3     OurList tail();
4 }
5 long sum(OurList xs) {
6     long total = 0;
7     for (; xs != null; xs = xs.tail())
8         total += xs.head();
9     return total;
10 }
```

Not only is this more reliable, but it's also quite a bit faster and follows normal Java idioms more closely. It also creates some opportunities to talk about some edge cases. For example, what if the list is dynamically generated and very long? The iterative solution is ideal here because it loses references to cells it's already processed, which allows them to be garbage collected. This means we can theoretically process lists of arbitrary length using constant space.

I also prefer the parameter name `xs` to a word like `list`. `xs`, which is pronounced, "exes," seems simpler to me than `list` does because it refers to a series of things called `x`, and `x` is the most generic name for a data value. It also takes up less space, which makes it easier to focus on the algorithm structure.

The `for` loop is a little nonstandard, but the idiom is preserved from Java's usual numeric iteration. This means that although someone unaccustomed to

this type of loop might be confused at first, the usual intuition about for loops would all apply: we're moving forward through a list until we hit the end, the termination check is the second thing, and the increment is the third.

5.2.3 Practice problems

Solutions to these are at [5.5](#).

1. $f(s) = \text{true}$ if s is a palindrome. Optimal solution is constant space, linear time, 6 lines total. Should take under a minute.
2. $f(xs) = \text{true}$ if xs contains a geometric series. Solution is constant space, linear time, 15 lines of code. Should take less than 5 minutes.
3. Write a function that counts the words in a string without allocating memory. Solution is linear time, 14 lines of code, and should take less than 5 minutes.

5.3 Algorithm problems

These usually have a known-best solution, but it's ok not to find it. You should of course try, and brush up on all of the usual algorithm and data structure stuff before a coding interview in case you get a problem like this. Below I go through some example problems. It's been long enough since I've seen these that I don't remember the solution.

By the way, you'll notice the almost complete absence of comments from my example code. Most interview problems are simple enough that comments are unnecessary. When you're practicing questions like this, write your code without comments or blank lines and see how readable you can make it. It's ok if you later decide to use comments, but most really good code won't require them.

5.3.1 Maximum subarray

Given a list of integers, find the offset and length of the subarray with the largest sum.

Initial train of thought It isn't immediately obvious how to do this, so start by talking through a few strategies. One option is to calculate the sum of every subarray and find the maximum of those, but it seems slow. These subarrays are always made from adjacent elements, though, so maybe we can ask something like, "should I expand out by one." In fact, since we're just summing stuff, we can expand out anytime we have something nonnegative and we'll get a larger, or at least not smaller, sum.

If I were to start coding now (which would be a mistake), here's what I would write:

Listing 5 examples/max-subarray1.java

```
1 // return [offset, length]
2 int[] subarray(final int[] xs) {
3     int nextUntried = 1;
4     int maxSum      = 0;
5     int maxSumStart = -1;
6     int maxSumLength = 0;
7     for (int i = 1; i < xs.length; i = nextUntried) {
8         maxSumStart = -1;
9         maxSumLength = 1;
10        while (i + maxSumStart >= 0 && xs[i + maxSumStart] >= 0) {
11            --maxSumStart;
12            ++maxSumLength;
13        }
14        i += maxSumStart;
15        while (i + maxSumLength < xs.length
16            && xs[i + maxSumLength] >= 0)
17            ++maxSumLength;
18        nextUntried = i + maxSumLength + 1;
19    }
20    return new int[] {maxSumStart, maxSumLength};
21 }
```

The code is a little awkward, but it isn't terrible. The problem is with my thinking: it's obvious I don't really understand something important about the problem. The right move is to continue thinking about it.

First course correction Ok, so what happens when we're expanding and hit a negative number? I guess we need to look on the other side to see if there's an even larger positive number. That's kind of gross, and bumps us back into quadratic time; so let's think about it differently. Suppose we first take a running sum of the array; then what are we looking for? Oh! We're just finding the max and min of that. The numbers in between those extremes collectively have the largest sum.

Writing code here is a much less serious mistake, but is still suboptimal. Here's what I'd say, most likely:

Listing 6 examples/max-subarray2.java

```
1 int[] subarray(int[] xs) {
2     // modify xs in place (!)
3     for (int i = 0, total = 0; i < xs.length; ++i)
4         xs[i] = total += xs[i];
5
6     int max = 0;
7     int min = 0;
8     for (int i = 1; i < xs.length; ++i) {
```

```

9      if (xs[i] > xs[max]) max = i;
10     if (xs[i] < xs[min]) min = i;
11 }
12 return new int[] {Math.min(min, max),
13                  Math.max(min, max) - Math.min(min, max)};
14 }

```

This is much better, but it's wrong and still suboptimal. Let's continue thinking:

Second course correction Ok, so how do we do this in code? We could start by summing the array in a loop, then check for max and min. Oh, but what if min happens after max? I guess we really need the largest min/max gap where the max is after the min.

Do we actually need to sum the array in a loop? We never look backwards at stuff we've summed, so really we just need to go through and maintain a running total. This is an optimal solution, so coding now makes sense:

Listing 7 examples/max-subarray3.java

```

1 int[] subarray(final int[] xs) {
2     if (xs.length == 0)
3         return null;
4     int max = 0;
5     int min = 0;
6     int maxValue = xs[0];
7     int minValue = xs[0];
8     for (int i = 1, sum = xs[0]; i < xs.length; ++i) {
9         sum += xs[i];
10        if (sum < minValue) {
11            minValue = maxValue = sum;
12            min = max = i;
13        }
14        if (sum > maxValue) {
15            maxValue = sum;
16            max = i;
17        }
18    }
19    return new int[] {min, max - min};
20 }

```

This solution feels a little more complicated than the last one, but it's conceptually more straightforward (since we're not modifying anything) and better-engineered. Just the fact that we're solving it in constant space is a consideration that a good interviewer would appreciate.

5.3.2 Parse simple math

*Write a function that takes a string containing a simple expression and returns the result. The expression will contain single-digit numbers, +, and *, and should respect precedence.*

Ok, this is obviously a shunting-yard style problem. We could parse it recursive-descent, but that's overkill. All operators are left-associative, so we can evaluate as we go. Actually, we can just have two constant levels: the number on the left of a + and the current product. So our solution is constant-space.

Listing 8 examples/parse-math.java

```
1 double parsemath(final String expr) {
2     double left = 0;
3     double right = 1;
4     for (int i = 0; i < expr.length(); ++i) {
5         final char c = expr.charAt(i);
6         switch (c) {
7             case '+':
8                 left += right;
9                 right = 1;
10                break;
11            case '*':
12                break;
13            default:
14                right *= c - '0';
15                break;
16        }
17    }
18    return left + right;
19 }
```

I had some doubts about 2/3 of the way through this. I wasn't sure it would be ok to add `right` to `left` at the end, since it might contain a stray 1 from either initialization or the + case. But after some thought I realized it would be ok since we always end after seeing a digit. That digit changes the 1 into something meaningful.

In hindsight, my variable names could be better. `left` and `right` are less descriptive than `sum` and `product`. Also, the first time through I had forgotten parens after `expr.length`. This isn't terrible, but it isn't great either. In my case I was confusing Java and Javascript, but an interviewer might assume I was bad at Java and confusing strings and arrays.

One nice part of this solution is `c - '0'`. Using the equivalent `c - 0x30` or `c - 48` would be hostile to someone reading your code, whereas `c - '0'` gives someone just enough hints that it's clear you're not up to anything too complicated.

Doing things like character-to-number conversion this way tells the interviewer you know how characters work and have at least some knowledge of ASCII. It's also reminiscent of C-style code, which will usually work in your favor in today's hiring culture of very high-level languages and app development (it shows perspective, assuming you can also do the high-level stuff).

5.3.3 Sorted word index

Build the index for a book, sorted by word. Your input is a list of $\langle \text{word}, \text{page} \rangle$ tuples, and the list may contain duplicates.

We'll need some data structures for this. I'll go ahead and write this up the wrong way to illustrate some common mistakes:

Listing 9 examples/sorted-index-bad.java

```
1 List index(List tuples) {
2     HashMap resultIndex = new HashMap();
3     for (int i = 0; i < tuples.size(); ++i) {
4         if (!resultIndex.containsKey(tuples[i].word))
5             resultIndex.put(tuples[i].word, new LinkedList());
6         resultIndex.get(tuples[i].word).add(tuples[i].page);
7     }
8     List sortedResult = new ArrayList();
9     Iterator iterator = resultIndex.iterator();
10    while (iterator.hasNext()) {
11        Tuple indexEntry = iterator.next();
12        sortedResult.add(indexEntry);
13    }
14    quicksort(sortedResult);
15    return sortedResult;
16 }
```

This solution has a lot of problems. For one thing, none of the data structures are parameterized. *Always parameterize your data structures.* There are two reasons you want to do this. First, prior to Java 1.5 (a long time ago), Java didn't even have generics. You might come across as a Luddite who prefers inferior technologies.

But more importantly, writing this stuff down tells the interviewer what you're thinking. A map of $\langle \text{Thing}, \text{Thing} \rangle$ doesn't convey your intentions very well; its meaning is then defined by what you do to the map instead of something intrinsic to it.³ Whenever possible, use code to document what you're doing. Especially parameters and return types, because this gives the interviewer a quick way to correct any inappropriate assumptions you might be making.

Here are some other problems with the code above:

³Generics aren't intrinsic really, but you get the idea.

1. It mixes definition types between interfaces (e.g. `List`) and implementations (e.g. `HashMap`). There are very occasionally reasons to do this, but this isn't one of them.
2. It grossly mishandles tuples. Were tuples a linked list, the function would be quadratic.
3. `resultIndex` is redundant, misleading, and too long. It takes up a huge amount of space inside the `for` loop.
4. `tuples[i].word` is repeated three times in quick succession. Stash this into a short variable.
5. `sortedResult` is too long.
6. `iterator` is redundant and too long. The ideal name is `it`.
7. The `while` loop spends a whole line defining a variable that's almost as long as the expression it replaces.
8. `quicksort()` is not only dangerous performance-wise, but also not a valid Java function. No Java library function would be called this way because Java requires that it be a method call (unless you `import static`, but your interviewer might not assume you meant that).

The code also has some technical issues that I don't think are important for an interview:

1. Maps don't support `.add`; the correct method is `.put`. Nobody cares. The compiler would catch this immediately. If the interviewer makes a big deal about this, find a company with better priorities or more competent interviewers.
2. `tuples[i]` doesn't work. This is fine; you clearly want to get-by-position. The only thing the interviewer should care about here is that this retrieval is linear-time for a `LinkedList`, and you should know this cold and have a very good excuse for doing it this way.
3. `.size()` could be expensive, but isn't. Your interviewer may not have read the source for `LinkedList`, but there's a counter in there that makes `.size()` a constant-time operation. It's cool if you can work this into an interview, but odds are that if you're counting a linked list you're already doing something wrong.

Here's the right way, with some notational liberties that are appropriate during an interview:

Listing 10 `examples/sorted-index-good.java`

```

1 SortedMap<String, List<Integer>> index(final List<Tuple> tuples) {
2   final SortedMap<...> result = new TreeMap<>();

```

```

3   for (final Tuple t : tuples) {
4       if (result.containsKey(t.word))
5           result.get(t.word).add(t.page);
6       else
7           result.put(t.word, new ArrayList<Integer>(t.page));
8   }
9   return result;
10 }

```

`SortedMap<...>` isn't allowed, but it's really obvious what I mean, especially if I explain it as I'm typing. The same thing is true of `new ArrayList<Integer>(t.page)`, which actually is valid Java but doesn't mean what I want it to (in real life it sets the initial capacity of the list). In an interview I would either make a point of mentioning this redefinition, or would just write out the right code.

It's worth mentioning that I'm being a little more pedantic than necessary about a few things. In particular, an interviewer is unlikely to care much whether you use `containsKey` vs `contains` and stuff like that. Knowing these details adds a bit to your credibility, and it's something I try to do, but it won't ever be the difference between passing and failing an interview.

5.4 High-level design questions

These are usually problems for which you don't write code, but instead just describe how you would design a solution. Failure and success are more nebulous here, but generally there are two ways to fail one of these questions:

1. The problem has a well-known solution and you completely miss it.
2. You fail to convince the interviewer that you could actually solve the problem given enough time.

An example of the first failure mode would be something like, "design a data structure that lets me get the minimum and maximum elements in $O(\log n)$ time" and you don't think of either a sorted list or a binary tree. You might be able to recover if you invent either one, but it's better (and far more reassuring to the interviewer) if you just know how they work.

The second failure mode happens either when you're completely out of your depth, or when you don't adopt the "give up or die trying" mentality thoroughly enough. If you're good at interviewing, neither should happen to you, and here's why. Being out of your depth isn't a dealbreaker if you're sufficiently resourceful and/or curious. If you don't have a clue, start asking definitional questions until the problem makes a little more sense. For example:

How many bits of entropy are in each word in English?

"Can I ask a question?"

"Sure, go for it."

"What is entropy?"

*"Entropy is information content, measured as Shannon Entropy."
"Do you mind if I look this up on Wikipedia?"*

You'll probably get mixed responses to the Wikipedia question. But it's a good question to ask because it shows that you know how to teach yourself to solve problems you can't yet work with. If the interviewer says yes you should be all set, but let's suppose they say no:

*"Nope, you've got to solve this problem with no external resources."
"Ok, I think I can solve this problem but not with just what I know. Can you relate entropy to other units of measure?"
"I guess you can think of entropy as the log of probability."
"Aha, so we're looking at words in terms of their probability distribution."
"That's right."
"Ok, so I guess more common words are less informative because they appear more often. Does this mean they have lower entropy?"*

This is naive and misguided question, but it forces the interviewer to either fix your thinking or be pathologically unhelpful. This is exactly the situation you want to create when you're operating at a disadvantage. Here's one way the conversation can proceed:

*"Well, a word by itself doesn't have entropy; it's a property of the probability distribution."
"Interesting, so we don't know the information content of a single word."
"That's right, because you need to consider its context to know how common it is."
"Ok, so entropy must have something to do with how the probabilities themselves are distributed."*

And so on. When the interviewer is your only source of information, it's critical to keep the conversation going and maintain enough curiosity to find your way around the problem space. Interviews are conversations as much as they are anything else, and how you think is as important as the knowledge you bring to the table.

It's worth noting that in order for this to work, you need to have a fairly broad base of knowledge to tie things to. The interviewer has neither the time nor inclination to teach you computer science or math from the ground up, so you have to be making some intelligent connections for them to be remotely confident that you'll work out.

5.4.1 IO scheduler

You're writing an operating system and inventing your own IO scheduling strategy. How would you design it?

This is a really open-ended question, but you'll get those sometimes. Here's how a good interview might go:

"Ok, so we're scheduling programs' requests to make IO?"
 "That's right."
 "Ok. These programs can read and write, and let's treat each request as having an offset and length from the device."
 "Yep, you can assume that."
 "I guess each program should have its own IO priority. So stuff like databases can get first access when it needs to."
 "Can you see anything that might go wrong if you do it this way?"
 "Well, theoretically we might swap to disk, and that's probably being scheduled too. But then, maybe we just say that it's inadvisable to deprioritize the swapper."
 "Fair enough."
 "Ok, so every process is going through the same scheduler. I guess the question now is what to do at any given moment. Let's keep a priority queue of requests."
 "How do you sort the queue?"
 "Program priority.... oh, and we also need to prevent starvation I guess. There's probably a standard way of dealing with this."
 "There is, but suppose you're inventing it from scratch."
 "Ok, well one way to do it is to keep track of the request times. As requests get older, they get bumped up the queue even if they're technically lower priority."
 "That sounds reasonable. What would that sorting function look like?"
 "I guess it would be something like $\text{priority} + k(\text{delay})$?"
 "How might that function not do what you want?"
 "I'm not sure. I guess if k is really small it would still allow starvation. But older requests will still outweigh new ones eventually."
 "Stepping back, what are some other considerations you might want to make in your scheduler?"
 "Let's see... we're handling priority correctly, so I'm not sure." "Ok, here's a hint. Suppose you've got two programs running at the same priority. One of them is writing 10GB in 4MB chunks, and the other one is writing all over the place in tiny 4KB chunks. What's going to happen here?"
 "Oh, that's awful. The random writer is going to introduce all kinds of seek delays into the sequential one. I should have thought of that."
 "Yep, exactly. What do you think is ideal behavior in this scenario?"
 "I'm not quite sure. Intuitively it seems like we should let the sequential one run uninterrupted for a while, but then what if you have two such processes? At that point you're balancing sequential access against delays."
 "Right, it's a tradeoff between latency and throughput. Is that a problem?"
 "I guess it isn't; that could be a system-wide preference."
 "Right. So any ideas about the sequential process?"

And so forth. Even though the interviewer had to prompt and give hints, it was mostly just to direct the conversation, not because the candidate was messing stuff up. Some stuff the candidate did right:

- Scoped out the problem before trying to solve it. The problem as stated actually wasn't solvable because it was underspecified, but the candidate asked enough questions to get something tractable.
- Recognized that real-world constraints applied to the problem and used them (e.g. mentioning disk swapping). Not all problems are amenable to this kind of reasoning, but this one is. An interviewer will usually appreciate this kind of perspective because junior developers often don't have it.
- Pointed out that the problem was most likely solved already. This is a good thing to do because the interviewer might immediately cut over to asking about established solutions. If you don't offer an existing solution, an interviewer might suspect you reinvent the wheel too eagerly.
- Described the solution at the same level of detail as the interviewer constrained the problem. The candidate didn't burst into code or mention direct memory access. Either of these would be addressed at a different level of problem, and the candidate's job is to recognize this.
- Offered some relevant technical details. Mentioning a priority queue isn't out of place at this point because priority queues are ideas, not data structures. Had the candidate mentioned using a minheap, the interviewer would have had leverage to change the problem description in a way that would require the candidate to back out of that recommendation (maybe the sorting criterion changes over time, so you'd have to re-heapify everything). This wastes time and suggests that the candidate might be prematurely optimizing or otherwise operating at the wrong level of abstraction.
- Remained willing to change the solution's design. Because the problem isn't written down anywhere, it's going to change constantly until the interview ends. It's important to treat it this way and not commit to anything too low-level. You want to remain as agile as the interviewer.

I tried to write an example of a bad interview, but I couldn't figure out how the conversation would actually go. A lot of times a bad interview is just someone who gets very quiet, or who doesn't understand the interviewer and doesn't then ask questions to fix it. A curious person who's sufficiently out of the loop might still fail a question like this, but odds of success go way up if you stay engaged.

5.5 Practice problem solutions

1. Return true if palindrome.

Listing 11 solutions/palindrome.java

```
1 boolean palindrome(final String s) {
```

```

2   for (int i = 0, j = s.length() - 1; i < j; ++i, --j)
3       if (s.charAt(i) != s.charAt(j))
4           return false;
5   return true;
6   }

```

2. Return true if geometric series.

Listing 12 solutions/geometric.java

```

1  boolean geometric(final double[] xs) {
2      if (xs.length < 2)
3          return true;
4      if (xs[0] == 0) {
5          for (int i = 0; i < xs.length; ++i)
6              if (xs[i] != 0)
7                  return false;
8          return true;
9      }
10
11     final double ratio = xs[1] / xs[0];
12     for (int i = 2; i < xs.length; ++i)
13         if (Math.abs(xs[i] / xs[i - 1] - ratio) > EPSILON)
14             return false;
15     return true;
16 }

```

3. Return the word count of a string without allocating memory.

Listing 13 solutions/wordcount.java

```

1  int wordcount(final String s) {
2      boolean inWord = false;
3      int words = 0;
4      for (int i = 0; i < s.length(); ++i) {
5          final char c = s.charAt(i);
6          if (c >= 'A' && c <= 'Z' || c >= 'a' && c <= 'z') {
7              if (!inWord)
8                  ++words;
9              inWord = true;
10         } else
11             inWord = false;
12     }
13     return words;
14 }

```